

Artificial Neural Networks (ANN) and K Nearest Neighbors (KNN)

AMIT V GOTTIPATI

Introduction

As part of this binary classification project Artificial Neural Networks (ANN) and K Nearest Neighbors were implemented in Python and Google colab using the SGEMM GPU kernel performance data set from UCI Machine Learning repository -> Dataset and on the Rain in Australia dataset from Kaggle -> Dataset2 . The project involved data pre-processing, data visualization, algorithm application, prediction and experimenting with different parameters like Number of hidden layers, Neurons, Activation Functions, Optimizer, Batch size, epoch, K, etc.

Dataset Description

Dataset1 - SGEMM GPU :

- The First dataset contains 241600 observations and 18 variables with no missing values.
- There are 18 variables, the first 14 are parameters upon which we will train our model. The first 10 parameters are ordinal with 4 levels, the next 4 are binary and last 4 variables are 4 different run times for the SGEMM GPU
- The description for the independent variables is given in the dataset link.
- The dataset was divided into train and test with a 70:30 split.
- For dependent variable 'avg' was created using average of the 4 run times provided which tells us the average running time for matrix-matrix product of the four runs. For classification the dependent variable 'avg' was converted into a binary variable using median as the threshold. (High/Low)

Dataset2:

- This Dataset contains 142193 observations with 24 variables
- Here the Target variable is Rain Tomorrow, we need to predict if it will rain tomorrow or not.
- It is given that Risk_MM variable is a very good predictor of the outcome and is dropped.
- Variables with more than 35% missing values were dropped and Last observation carried forward LOCF was used to impute the remaining missing values

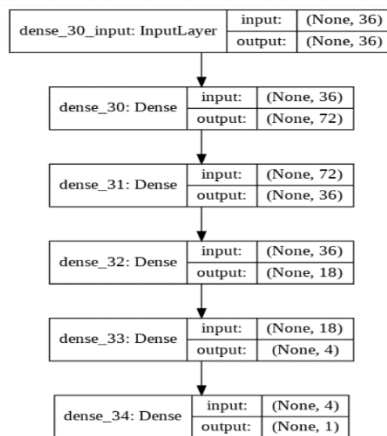

Artificial Neural Networks

Dataset-1:


Tanh = T	Sigmoid = S	Relu = R	Adam = A	SGD =SG	SoftMax = SX
-------------	----------------	-------------	-------------	------------	-----------------

Model #	No of Hidden Layers	Neurons in each Hidden layer	Activation Function in Each Layer (Hidden & Output Layer)	Optimizer	Batch Size	Epoch	Loss	Train Error	Test Error	Training Time
1	2	4,4	R,R,S	A	10	5	0.1221	0.05162	0.052	207
2	2	4,4	R,R,S	A	20	5	0.1231	0.0508	0.0515	130
3	2	4,4	R,R,S	A	100	5	0.116	0.0492	0.0498	30
4	2	4,4	R,R,S	A	10	10	0.1094	0.0456	0.0456	517
5	4	4,4,4,4	R,R,R,R,S	A	10	5	0.6934	0.4999	0.5005	318
6	4	4,4,4,4	R,R,R,R,S	A	100	10	0.6931	0.4999	0.5005	70
7	3	4,4,4	R,R,R,S	A	10	5	0.1838	0.0689	0.0674	278
8	3	4,4,4	R,R,R,S	A	50	5	0.174	0.0651	0.0647	61
9	3	4,4,4	R,R,T,S	A	10	5	0.1219	0.0514	0.0509	292
10	3	4,4,4	R,R,T,S	SG	10	5	0.1242	0.0513	0.0509	216
11	2	4,2	R,T,S	SG	10	5	0.1397	0.0581	0.0584	205
12	2	4,2	R,T,S	A	10	5	0.1241	0.0522	0.0542	269
13	4	72,36,18,4	R,T,T,T,S	A	10	5	0.0469	0.0181	0.0201	289
14	4	72,36,18,4	R,T,R,T,S	A	10	5	0.0509	0.0208	0.0221	292
15	5	36,36,18,18,4,1	R,T,S,S,T,S	A	10	5	0.0539	0.0218	0.023	303
16	9	36,32,28,24,20,16,12,8,4	R,T,T,T,T,T,S,S,T,S	A	10	5	0.0602	0.0245	0.0261	401

Training Time vs Training Error



Hidden Layers vs Error

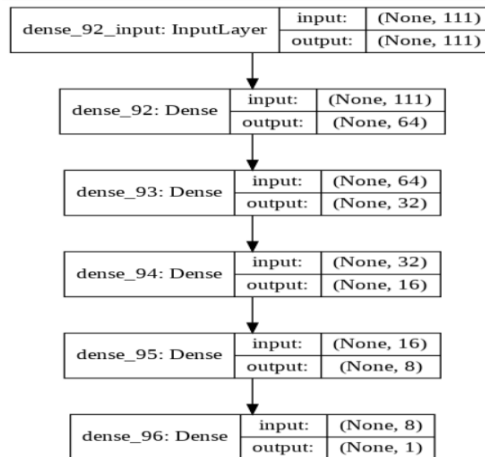



From this experiment we saw that Adam (RMSProp + Momentum) and SGD(Stochastic Gradient Descent) optimizer gave similar results. We got the best result when we used a 4 hidden layer neural network with Relu and Tanh activation functions with decreasing number of neurons. We also observed lowest loss for this result. Increasing the complexity of neural network by increasing number of hidden layers decrease the train and test error but it started increasing again with high number of hidden layers .


Dataset-2:

Dataset2										
Model #	No of Hidden Layers	Neurons in each Hidden layer	Activation Function in Each Layer (Hidden & Output Layer)	Optimizer	Batch Size	Epoch	Loss	Train Error	Test Error	Training Time
1	2	4,4	R,R,S	A	10	5	0.3455	0.1482	0.1516	70
2	2	4,4	R,R,S	A	20	5	0.3426	0.1473	0.1489	40
3	2	4,4	R,R,S	A	100	5	0.3398	0.146	0.1484	20
4	4	4,4,4,4	R,R,R,R,S	A	10	5	0.3526	0.1495	0.1532	78
5	4	64,32,16,8	R,R,R,R,S	A	10	5	0.3092	0.1309	0.1429	83
6	4	64,32,16,8	R,T,T,R,S	A	10	5	0.3064	0.1308	0.1418	82
7	4	64,32,16,8	R,S,S,R,S	A	10	5	0.3102	0.1328	0.1409	82
8	4	64,32,16,8	R,S,S,R,S	SG	10	5	0.2983	0.127	0.1399	75
9	4	64,32,16,8	R,S,S,R,SX	SG	10	20	11.84	0.7771	0.7731	320
10	4	64,32,16,8	T,T,T,T,S	SG	10	20	0.3309	0.14	0.1455	300
11	4	64,64,64,64	T,T,T,T,S	SG	10	5	0.3581	0.1517	0.1544	78
12	1	64	T,S	SG	10	5	0.3571	0.1519	0.1562	70
13	11	100,90,80,70,60,50,40,30,20,10,5	T,T,T,T,T,S,S,T,T,T,T,S	SG	10	5	0.5307	0.223	0.223	95
14	7	200,100,100,100,75,50,25	T,T,T,T,T,T,T,S	SG	10	5	0.3617	0.1527	0.1573	96

Training Time vs Training Error



Hidden Layers vs Error



From this experiment we got the best result when we used a 4 Hidden layer neural network with decreasing number of neurons, Relu and Sigmoid Activation functions and SG optimizer. The Training and Testing Error remains very similar for different neural networks.

K Nearest Neighbors (KNN)

We have used PyCaret library in python to implement KNN. For initial models we have used Minkowski metric for Dataset1 and Dataset2. Minkowski distance is a generalization of both the Euclidean distance and the Manhattan distance and hence can be a good fit for our model. Using Pycaret we have created tuned and calibrated models for KNN. We have used a 2 fold cross validation for all the models.

Dataset-1:

Train and Test Errors Respectively

1)KNN

	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.9499	0.9863	0.9381	0.9608	0.9493	0.8998
1	0.9494	0.9862	0.9367	0.9612	0.9488	0.8989
Mean	0.9497	0.9862	0.9374	0.9610	0.9490	0.8994
SD	0.0002	0.0001	0.0007	0.0002	0.0003	0.0005

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 K Neighbors Classifier	0.9655	0.9932	0.9517	0.9787	0.965	0.931

2)Tuned KNN

	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.9594	0.9935	0.9401	0.9779	0.9586	0.9188
1	0.9576	0.9931	0.9388	0.9754	0.9568	0.9152
Mean	0.9585	0.9933	0.9394	0.9766	0.9577	0.9170
SD	0.0009	0.0002	0.0006	0.0012	0.0009	0.0018

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 K Neighbors Classifier	0.969	0.9965	0.9495	0.9879	0.9683	0.9379

4) Calibrated Isotonic KNN

	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.9511	0.9905	0.9254	0.9756	0.9498	0.9023
1	0.9500	0.9904	0.9239	0.9747	0.9486	0.8999
Mean	0.9505	0.9904	0.9247	0.9751	0.9492	0.9011
SD	0.0006	0.0001	0.0008	0.0004	0.0006	0.0012

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 Calibrated Classifier C V	0.9651	0.9943	0.9633	0.9668	0.965	0.9302

```
1 print(knn)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=5, p=2,
weights='uniform')
```

```
1 print(tuned_knn)
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='manhattan',
metric_params=None, n_jobs=None, n_neighbors=14, p=2,
weights='distance')
```

```
1 print(calibrated_knn)
CalibratedClassifierCV(base_estimator=KNeighborsClassifier(algorithm='auto',
leaf_size=30,
metric='minkowski',
metric_params=None,
n_jobs=None,
n_neighbors=5, p=2,
weights='uniform'),
cv=2, method='sigmoid')
```


```
1 print(calibrated_knn_isotonic)
CalibratedClassifierCV(base_estimator=KNeighborsClassifier(algorithm='auto',
leaf_size=30,
metric='minkowski',
metric_params=None,
n_jobs=None,
n_neighbors=5, p=2,
weights='uniform'),
cv=2, method='isotonic')
```

3)Calibrated KNN (Platt sigmoid)


	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.9551	0.9905	0.9552	0.9550	0.9551	0.9102
1	0.9529	0.9903	0.9539	0.9519	0.9529	0.9057
Mean	0.9540	0.9904	0.9546	0.9535	0.9540	0.9080
SD	0.0011	0.0001	0.0006	0.0015	0.0011	0.0022

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 Calibrated Classifier C V	0.9651	0.9943	0.9633	0.9668	0.965	0.9302

1)





2)



3)

4)

We

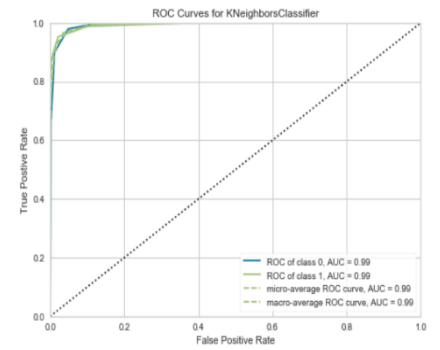
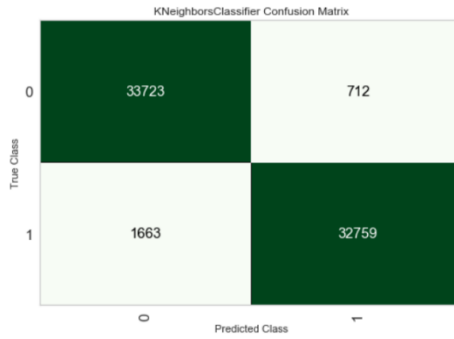
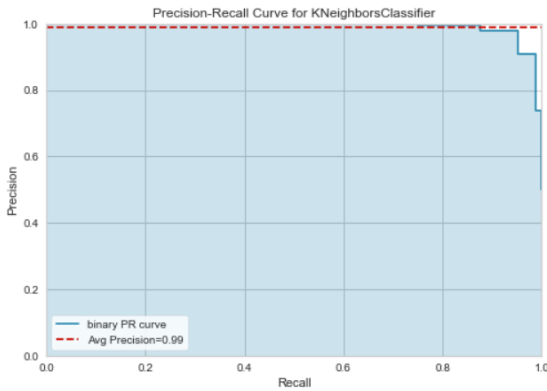


observe


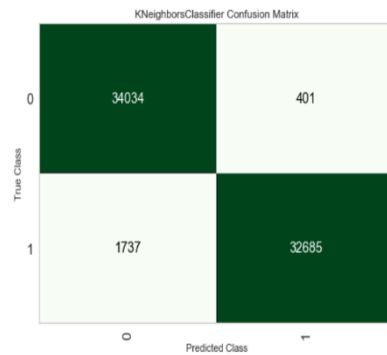

experiment that Tuned KNN gave us the best results with K=14 and manhattan as the distance metric. We also see that calibration plots for KNN and Tuned KNN are s curved. We tried to calibrate them using Platt sigmoid curve and isotonic regression and gave similar results.

from this

1)KNN



2)Tuned KNN



We ran all model comparison to see which model performed best. We can see that CatBoostClassifier gives best training accuracy. We need to apply on test data to see if overfitting is not taking place. A 2 fold cross validation has been applied to all models.

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	CatBoost Classifier	0.9839	0.999	0.9785	0.9893	0.9838	0.9678
1	Decision Tree Classifier	0.9834	0.9834	0.9841	0.9827	0.9834	0.9668
2	Extra Trees Classifier	0.9823	0.9984	0.9757	0.9888	0.9822	0.9646
3	Random Forest Classifier	0.9722	0.996	0.9571	0.9869	0.9718	0.9444
4	Light Gradient Boosting Machine	0.9672	0.996	0.9485	0.9853	0.9666	0.9344
5	K Neighbors Classifier	0.9497	0.9862	0.9374	0.961	0.949	0.8994
6	Gradient Boosting Classifier	0.9122	0.9614	0.8919	0.9297	0.9104	0.8244
7	Extreme Gradient Boosting	0.908	0.9604	0.8895	0.9236	0.9062	0.816
8	SVM - Linear Kernel	0.839	0	0.8359	0.841	0.8385	0.6779
9	Logistic Regression	0.8227	0.8842	0.8211	0.8237	0.8224	0.6454
10	Ridge Classifier	0.8226	0	0.8155	0.8271	0.8213	0.6451
11	Linear Discriminant Analysis	0.8226	0.8832	0.8156	0.8271	0.8213	0.6451
12	Ada Boost Classifier	0.8005	0.8765	0.8202	0.7891	0.8044	0.6011
13	Naive Bayes	0.7876	0.8244	0.7371	0.8198	0.7763	0.5752
14	Quadratic Discriminant Analysis	0.6809	0.7299	0.711	0.6706	0.6902	0.3617

Dataset-2:

1)KNN

	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.8352	0.8181	0.4565	0.7053	0.5543	0.4588
1	0.8331	0.8175	0.4482	0.7006	0.5467	0.4504
Mean	0.8342	0.8178	0.4523	0.7030	0.5505	0.4546
SD	0.0010	0.0003	0.0041	0.0024	0.0038	0.0042


```
1 print(knn)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=5, p=2,
weights='uniform')
```

```
1 print(tuned_knn)
```



```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
metric_params=None, n_jobs=None, n_neighbors=10, p=2,
weights='distance')
```

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 K Neighbors Classifier	0.84	0.8314	0.478	0.7149	0.573	0.4793





KNeighborsClassifier Confusion Matrix

	0	1
0	29693	1734
1	4749	4349





2) Tuned KNN



	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	0.8417	0.8481	0.4384	0.7535	0.5543	0.4661
1	0.8418	0.8514	0.4380	0.7541	0.5542	0.4661
Mean	0.8418	0.8497	0.4382	0.7538	0.5542	0.4661
SD	0.0000	0.0016	0.0002	0.0003	0.0001	0.0000

Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0 K Neighbors Classifier	0.846	0.8605	0.4647	0.7553	0.5754	0.4878

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa
0	CatBoost Classifier	0.8573	0.8848	0.5271	0.764	0.6238	0.5394
1	Light Gradient Boosting Machine	0.853	0.8768	0.5168	0.7509	0.6122	0.5253
2	Extra Trees Classifier	0.846	0.8709	0.4256	0.7922	0.5537	0.4707
3	Gradient Boosting Classifier	0.8459	0.8612	0.4782	0.7443	0.5822	0.4932
4	Extreme Gradient Boosting	0.8459	0.861	0.4681	0.7515	0.5769	0.4888
5	Logistic Regression	0.8439	0.8578	0.4908	0.7252	0.5854	0.4936
6	Linear Discriminant Analysis	0.842	0.8559	0.501	0.71	0.5875	0.4934
7	Ridge Classifier	0.8419	0	0.4396	0.7532	0.5552	0.467
8	SVM - Linear Kernel	0.8407	0	0.4573	0.7331	0.5631	0.4721
9	Ada Boost Classifier	0.8383	0.8511	0.474	0.7091	0.5682	0.4735
10	Random Forest Classifier	0.8353	0.8309	0.4122	0.7385	0.529	0.4388
11	K Neighbors Classifier	0.8342	0.8178	0.4523	0.703	0.5505	0.4546
12	Decision Tree Classifier	0.7822	0.6882	0.5175	0.5148	0.5162	0.3756
13	Naive Bayes	0.6258	0.7112	0.6883	0.3369	0.4523	0.216
14	Quadratic Discriminant Analysis	0.5817	0.6343	0.6291	0.2959	0.4009	0.1396



We can see from this experiment that on dataset 2 the Tuned KNN performed better as compared to KNN. Both are very well calibrated and so we don't need to do any calibration. The Tuned KNN used Euclidean distance with K=10. We also KNN performance when compared to other algorithms and again see that CatBoost performs Best.

Comparison of Both Algorithms:

Dataset 1 - ANN Performed Best, Accuracy was the Measure and it is a balanced dataset.

Dataset 2 - ANN Performed Best,

KNN took lot of time to train as compared to ANN, which provided higher accuracy in lesser time. This is because KNN takes more time when features are more. So when we dummy our variables the cost of computation is higher for KNN.

We can improve performance in Dataset2 by using synthetic sampling like SMOTE, ADASYN as it is imbalanced.